


# Deadlocks

# The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process
- Example1
  - System has 2 tape drives
  - $P_1$  and  $P_2$  each hold one tape drive and each needs another one
- Example2
  - Semaphores  $A$  and  $B$ , initialized to 1

$P_0$	$P_1$
Wait (A);	
	Wait(B)
	Wait(A)
Wait (B);	

# Bridge Crossing Example



- Traffic only in one direction.
- Each section of a bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation is possible.

# System Model

- Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances
- Each process utilizing a resource:
  - Makes Request for resource
  - Uses the resource
  - Releases the resource
- Request and release of resources is done through system calls

# Deadlock Characterization

Deadlock can arise if following **four** conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource(i.e. resource is non-sharable).
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that

$P_0$  is waiting for a resource that is held by  $P_1$

$P_1$  is waiting for a resource that is held by  $P_2$

$\dots$ ,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$

and  $P_n$  is waiting for a resource that is held by  $P_0$ .


# Resource-allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .

- $V$  is partitioned into two types:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system.
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- Request edge – directed edge  $P_i \rightarrow R_j$
- Assignment edge – directed edge  $R_j \rightarrow P_i$

# Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances



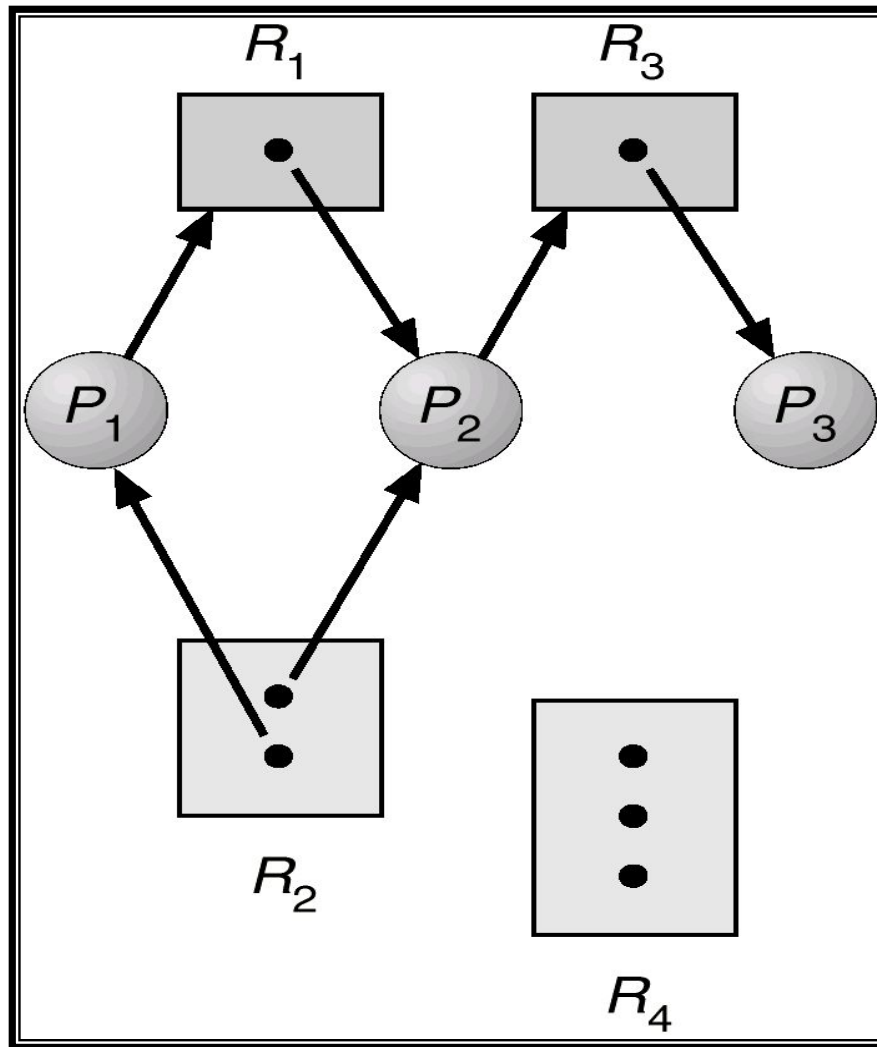
- $P_i$  requests instance of  $R_j$



- $P_i$  is holding an instance of  $R_j$

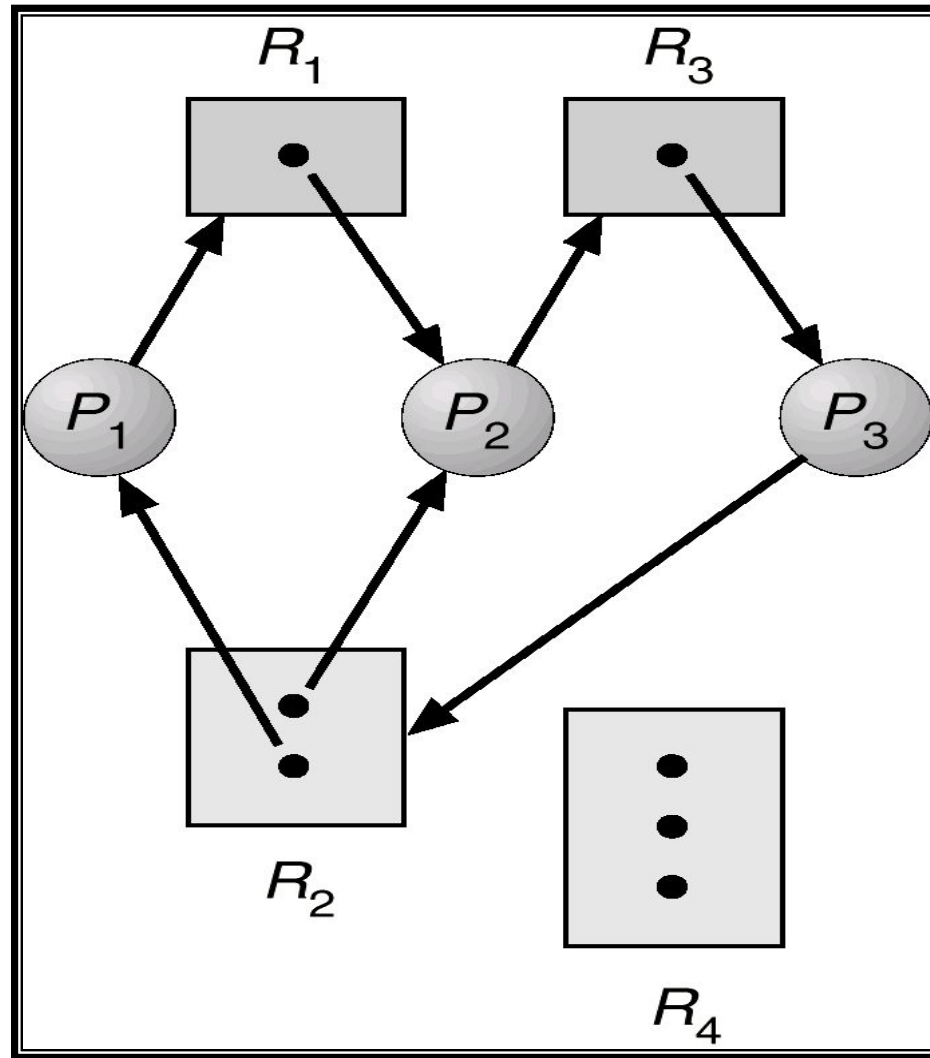


## Example of a Resource Allocation Graph

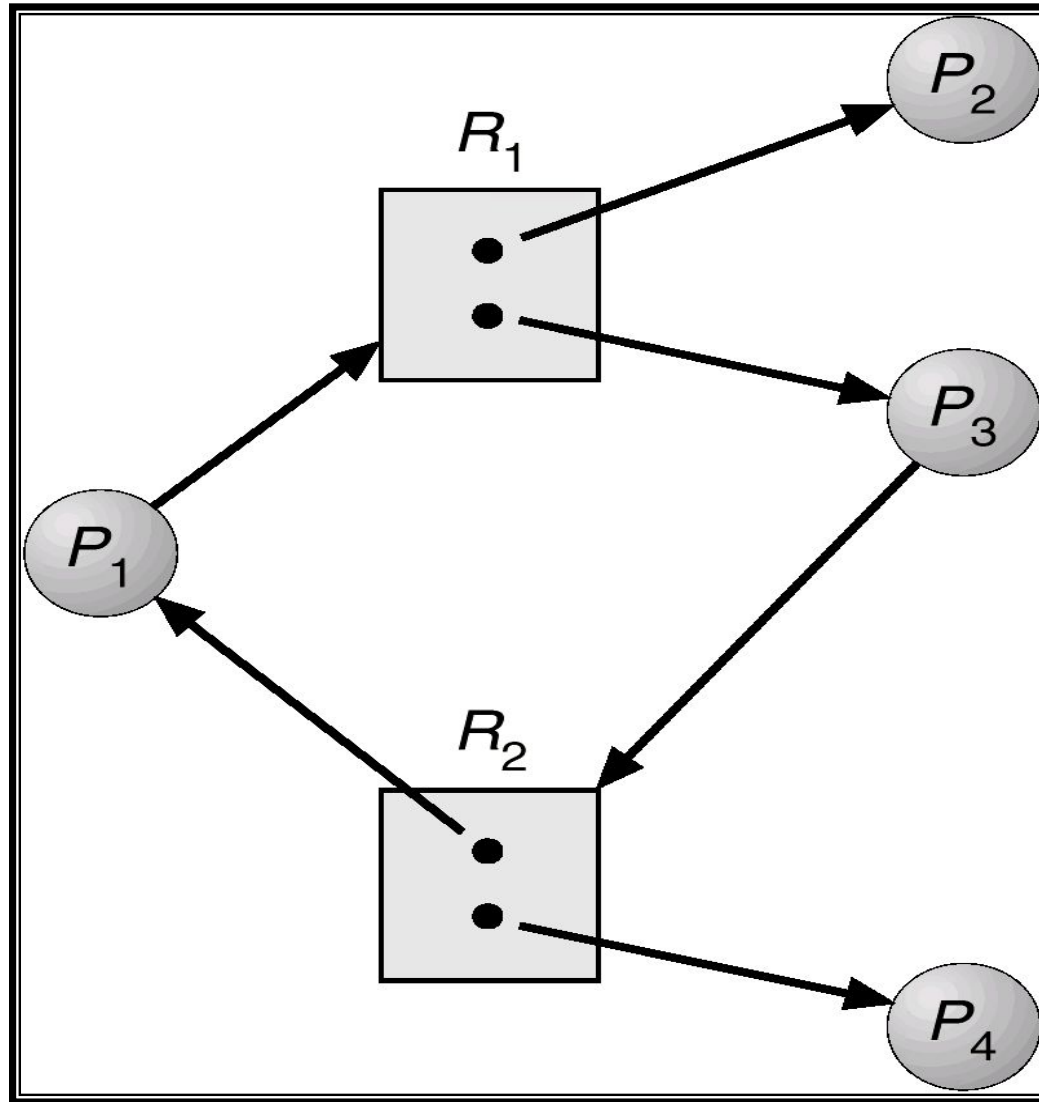




# Resource Allocation Graph With A Deadlock



## Resource Allocation Graph With A Cycle But No Deadlock



# Basic Facts

- If graph contains no cycles  $\Rightarrow$  no deadlock.
- If graph contains a cycle  $\Rightarrow$ 
  - ☐ If only one instance per resource type, then deadlock
  - ☐ If several instances per resource type, possibility of deadlock

# Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state.  
(Prevention or Avoidance)
- Allow the system to enter a deadlock state and then Detect & recover.

# Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion**
- **Hold and wait**
- **No preemption**
- **Circular wait**

# Deadlock Prevention

Ensure that one of the 4 necessary conditions for deadlock does not hold.

Constraints the way requests can be made for resources

- **Mutual exclusion**

- not required for sharable resources;
- Must hold for non-sharable resources.

- ☐ Make non-sharable resources as sharable

- ☐ some resources are intrinsically non sharable

# Deadlock Prevention

- **Hold and wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - Require process to request and be allocated all its resources before it begins execution,
    - Need to know all the requirement in the beginning itself
  - Allow process to request resources only when the process has none.
    - If a process has resource and requires additional resources, it must release all resources it is holding and then make request
  - Low resource utilization; Starvation possible

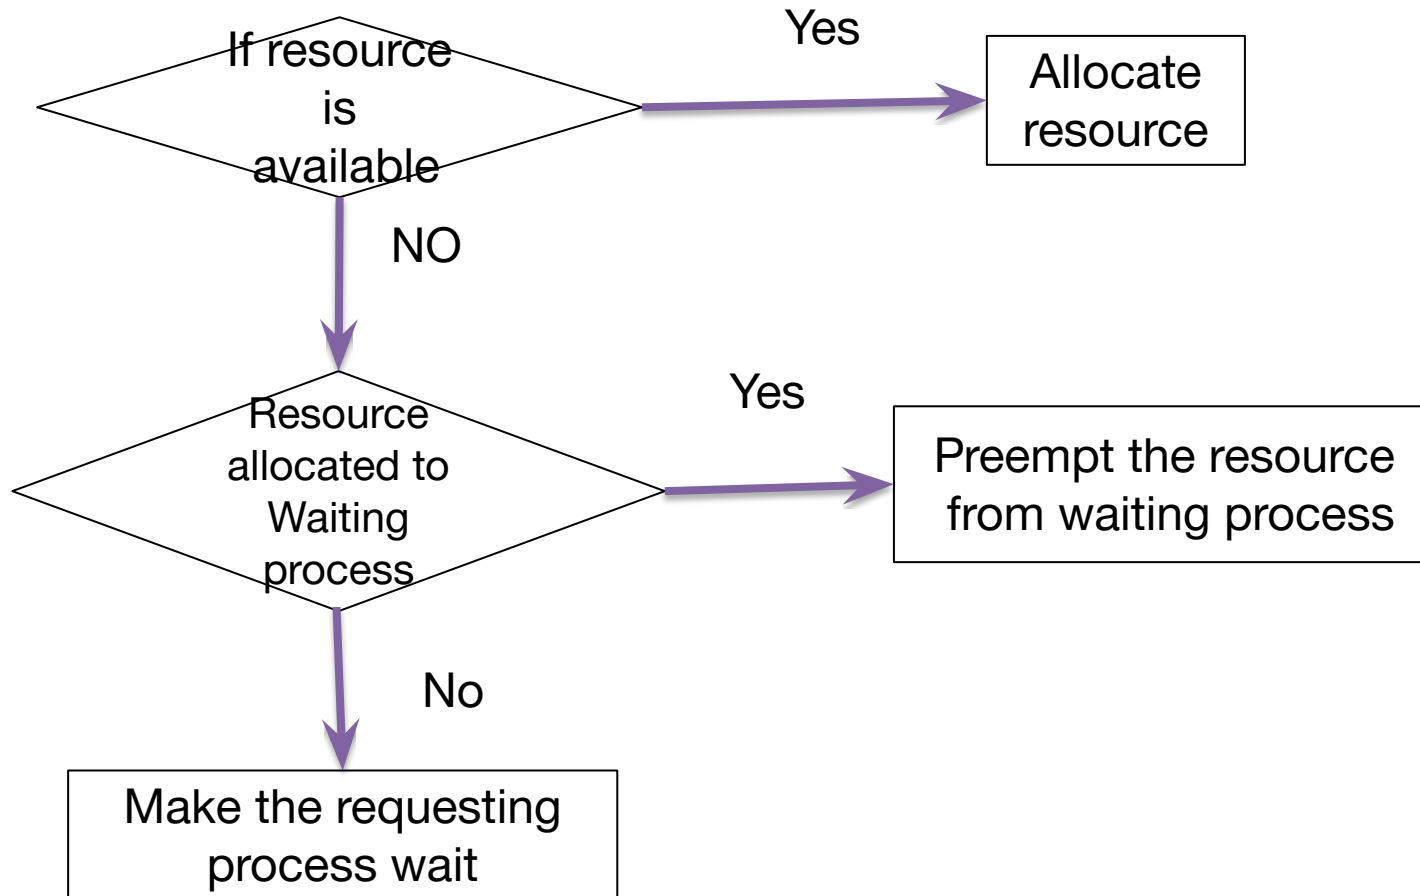
# Deadlock Prevention (Cont.)

- **No preemption – Approach -1**

- ☐ If a process that is holding some resources requests another resource that cannot be immediately allocated to it,
  - Release all resources currently being held by process
- ☐ Preempted resources are added to the list of resources for which the process is waiting
- ☐ Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.
- ☐ Cannot be generally applied to resources like printers, tape drives etc.



# No preemption (cont-)



# Deadlock Prevention (Cont.)

- **Circular wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

# Deadlock Avoidance

Requires that the system has some additional *a priori* information available.

- Simplest model requires that each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm
  - Dynamically finds the *resource-allocation state* in order to *ensure that a circular-wait condition does not occur*.
- Resource-allocation *state* is defined by
  - the number of available resources
  - the number of allocated resources,
  - the maximum requirement of the processes

# Safe State

- When a process requests an available resource, the system needs to decide if immediate allocation leaves the system in a *safe state*.
- A state is safe if system can allocate resources to each process (Max requirement ) in some order and still avoid deadlock
- System is in safe state if there exists a *safe sequence* of all processes.

- Sequence  $\langle P_1, P_2, \dots, P_n \rangle$  is safe :
- if for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$ 
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on

# Example


	Maximum need	Allocation
P0	10	5
P1	4	2
P2	9	2

Total no of resources =12  
safe sequence < P1,P0,P2>

# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks.
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock.
- Avoidance  $\Rightarrow$  ensure that a system will never enter into an unsafe state

# Safe, Unsafe , Deadlock State







# Resource-Allocation Graph Algorithm

- *Claim edge*  $P_i \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$ ; represented by a dashed line.
- Claim edge converts to *request edge* when a process requests a resource.
- When a resource is released by a process, *assignment edge* reconverts to a claim edge.
- Resources must be claimed *a priori* in the system.

# Resource-Allocation Graph For Deadlock Avoidance



# Unsafe State In Resource-Allocation Graph



# Banker's Algorithm

- Multiple instances.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.

## Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- **Available:** Vector of length  $m$ . If  $\text{Available}[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
- **Max:**  $n \times m$  matrix. If  $\text{Max}[i, j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
- **Allocation:**  $n \times m$  matrix. If  $\text{Allocation}[i, j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .
- **Need:**  $n \times m$  matrix. If  $\text{Need}[i, j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.

$$\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$$

# Safety Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively.  
Initialize:  
$$Work = Available$$
$$Finish[i] = false \text{ for } i = 0, 1, \dots, n-1.$$
2. Find an index *i* such that both:
  - (a) *Finish* [*i*] == *false*
  - (b)  $Need_i \leq Work$If no such *i* exists, go to step 4.
3.  $Work = Work + Allocation_i$   
*Finish*[ *i* ] = *true*  
go to step 2.
4. If *Finish* [*i*] == *true* for all *i*, then the system is in a safe state.

# Safe state example

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim Matrix

	R1	R2	R3
P1	1	0	0
P2	6	1	2
P3	2	1	1
P4	0	0	2

Allocation Matrix

R1	R2	R3
9	3	6

Resource Vector

R1	R2	R3
0	1	1

Available Vector

**(a) Initial state**

	R1	R2	R3
P1	2	2	2
P2	0	0	1
P3	1	0	3
P4	4	2	0

# Resource-Request Algorithm for Process $P_i$

$Request_i$  = request vector for process  $P_i$ .

If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$ .

1. If  $Request_i \leq Need_i$  then go to step 2.  
Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If  $Request_i \leq Available$ , go to step 3.  
Otherwise  $P_i$  must wait, since resources are not available.
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:  
 $Available = Available - Request_i;$   
 $Allocation_i = Allocation_i + Request_i;$   
 $Need_i = Need_i - Request_i$ 
  - If safe  $\Rightarrow$  the resources are allocated to  $P_i$ .
  - If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored



# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$
- 3 resource types  $A$  (10 instances),  $B$  (5 instances) and  $C$  (7 instances).
- Snapshot at time  $T_0$ :

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	$A$	$B$	$C$	$A$	$B$	$C$	$A$	$B$	$C$
$P_0$	0	1	0	7	5	3	3	3	2
$P_1$	2	0	0	3	2	2			
$P_2$	3	0	2	9	0	2			
$P_3$	2	1	1	2	2	2			
$P_4$	0	0	2	4	3	3			

# Example (Cont.)

- The content of the matrix Need is equal to Max – Allocation.

Need

*A B C*

$P_0$  7 4 3

$P_1$  1 2 2

$P_2$  6 0 0

$P_3$  0 1 1

$P_4$  4 3 1

	<u>Allocation</u>	<u>Max Available</u>		
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>	
$P_0$	0 1 0	7 5 3	3 3 2	
$P_1$	2 0 0	3 2 2		
$P_2$	3 0 2	9 0 2		
$P_3$	2 1 1	2 2 2		
$P_4$	0 0 2	4 3 3		

- The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria.

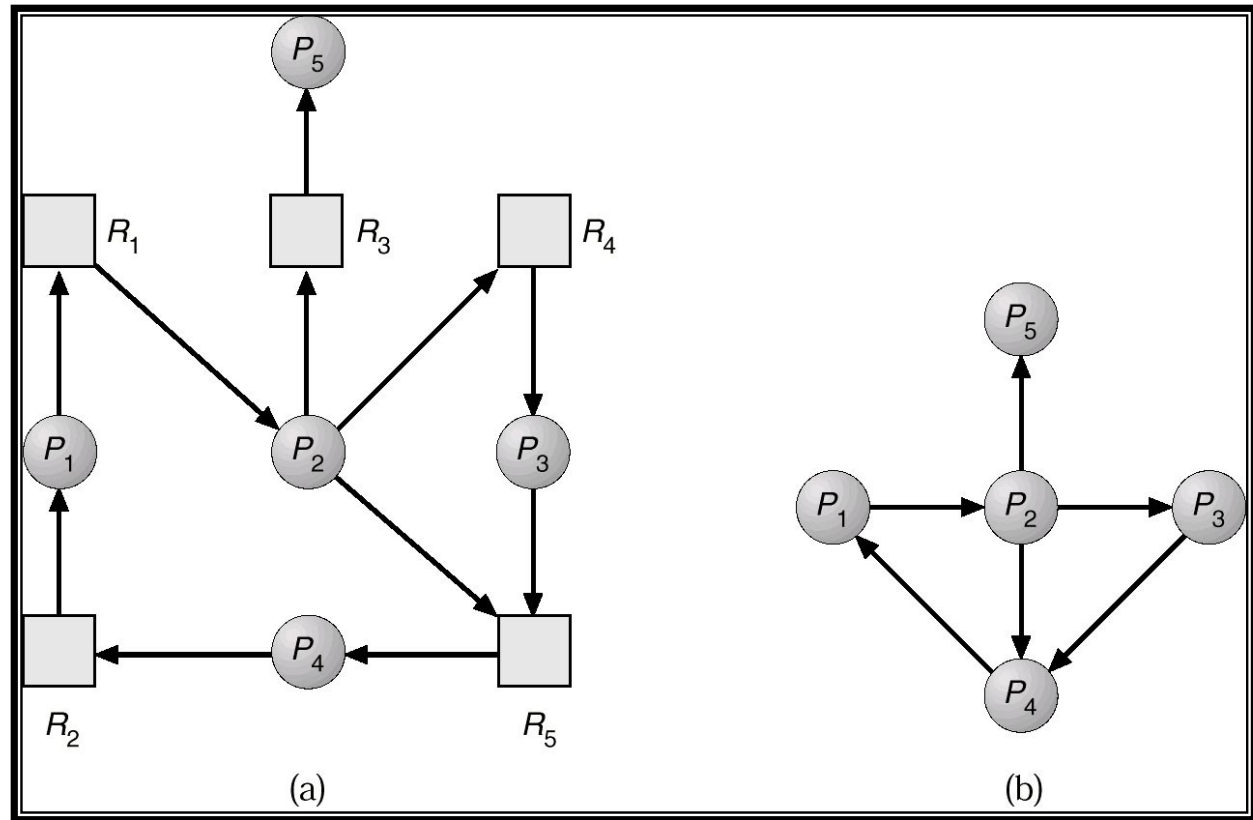
# Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

# Single Instance of Each Resource Type

- Maintain *wait-for* graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph.

# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation  
Graph

Corresponding wait-for  
graph

# Several Instances of a Resource Type

- *Available*: A vector of length  $m$  indicates the number of available resources of each type.
- *Allocation*: An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- *Request*: An  $n \times m$  matrix indicates the current request of each process. If  $Request[i, j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Detection Algorithm

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively. Initialize:

(a)  $Work = Available$

(b) For  $i = 1, 2, \dots, n$ , if  $Allocation_i \neq 0$ , then  
 $Finish[i] = false$ ; otherwise,  $Finish[i] = true$ .

2. Find an index  $i$  such that both:

(a)  $Finish[i] == false$

(b)  $Request_i \leq Work$   
If no such  $i$  exists, go to step 4.

3.  $Work = Work + Allocation_i$   
 $Finish[i] = true$   
go to step 2.

4. If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked.

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances).

- Snapshot at time  $T_0$ :

	<i>AllocationRequest Available</i>								
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	0			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

- Sequence  $\langle P_0, P_2, P_3, P_4, P_1 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$ .



## Example (Cont.)

- $P_2$  requests an additional instance of type C.

Request

A B C

$P_0$  0 0 0

$P_1$  2 0 1

$P_2$  0 0 1

$P_3$  1 0 0

$P_4$  0 0 2

- State of system?

- ☐ Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes requests.
- ☐ Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$ .

# Detection-Algorithm Usage

- When, and how often, to invoke depends on:
  - ☐ How often a deadlock is likely to occur?
  - ☐ How many processes will need to be rolled back?
- Frequency of execution
  - Invoke the algorithm every time request for resource cannot be granted immediately.
  - Once per hour
  - When CPU utilization drops below 40%

## Recovery from Deadlock: Process Termination

- Abort all deadlocked processes ☐ May be expensive
- Abort one process at a time until the deadlock cycle is eliminated ☐ Invoke deadlock detection algorithm after every abort.
- In which order should we choose to abort?
  - ☐ Priority of the process.
  - ☐ How long process has computed, and how much longer to completion.
  - ☐ Resources the process has used.
  - ☐ Resources process needs to complete.
  - ☐ Is process interactive or batch?

## Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.